실시간 빅데이터 처리 spark/flink
- Part5 에어플로우를 통한 배치 프로세싱

이 수업에서 다루고자 하는 것

이론

- Task & DAG
- Backfill & Catchup
- Timeout & Callback
- Retry & Alret
- Pool & Paralleism
- Hook
- SubDAG & TaskGroups
- Branching
- Trigger Rule
- XCOM

이 수업에서 다루고자 하는 것

실습

- PythonOperator
- BashOperator
- PostgresOperator
- BranchOperator
- BranchDateTimeOperator
- SubDagOperator
- TriggerDagRunOperator
- ExternalTaskSensor
- ShortCircuit Operator
- LatestOnlyOperator

Chapter 1. Abstraction

Chapter 1.

01. Architecture

01. Overview

- Workflow == DAG(Directed Acyclic Graph)
  - Contains individual pieces of work == Tasks

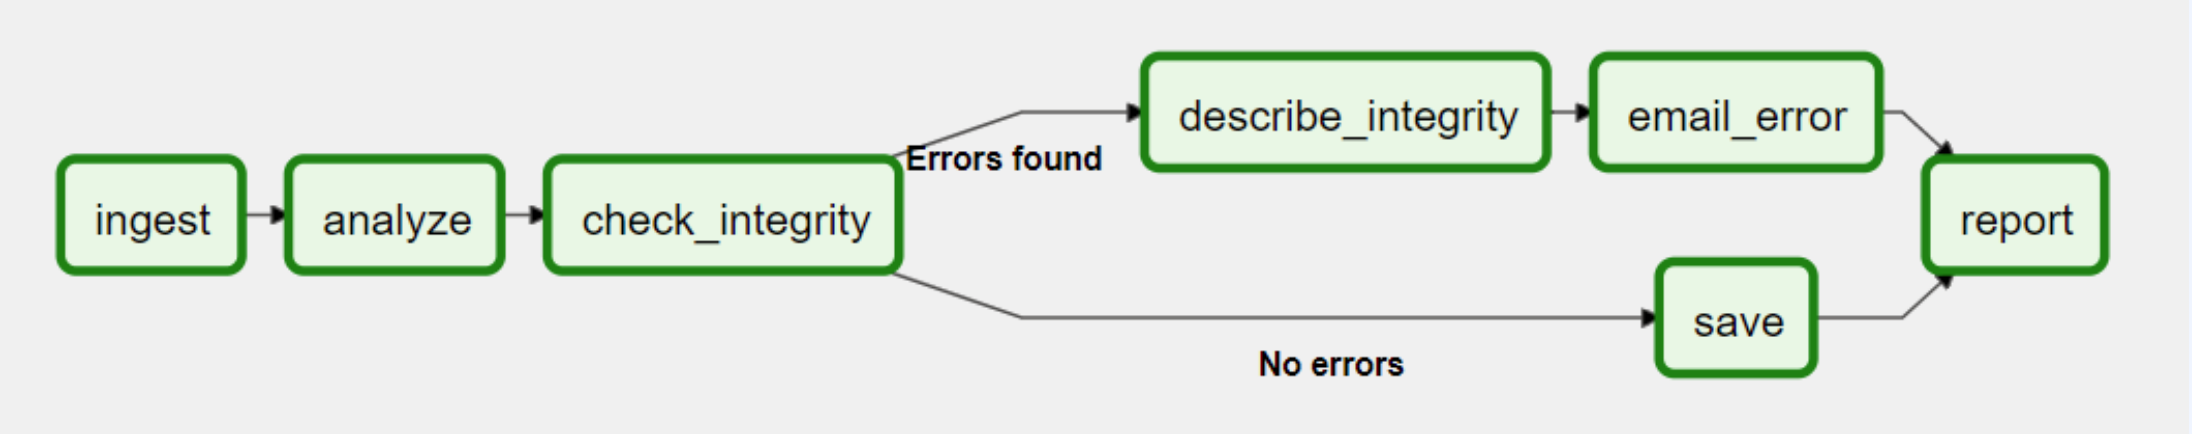

01. Components

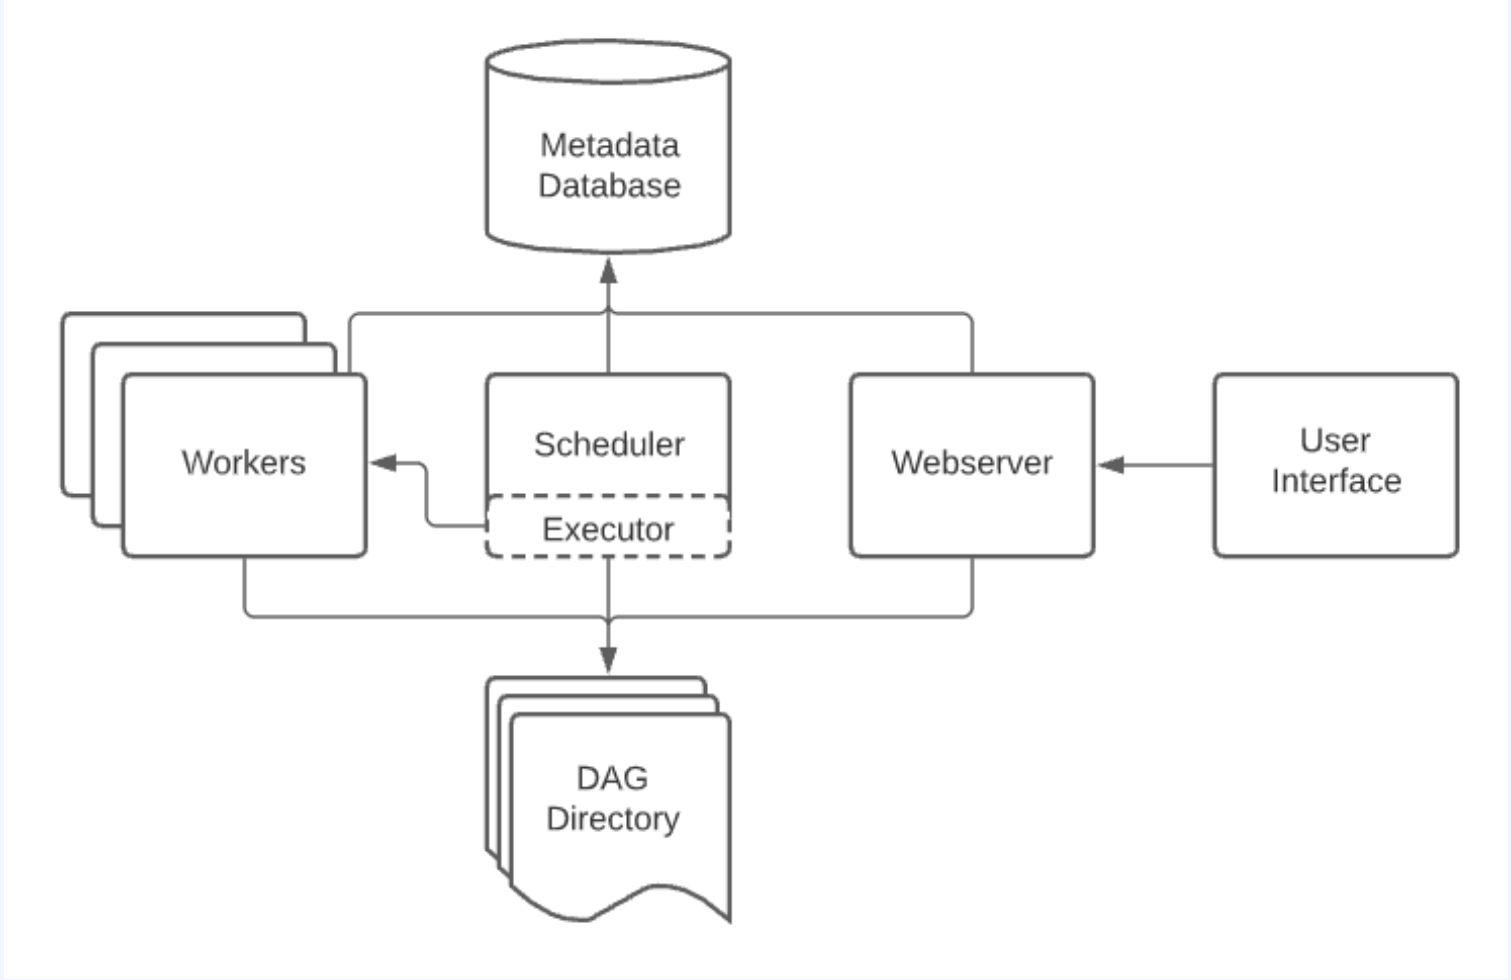

01. Task Lifecycle

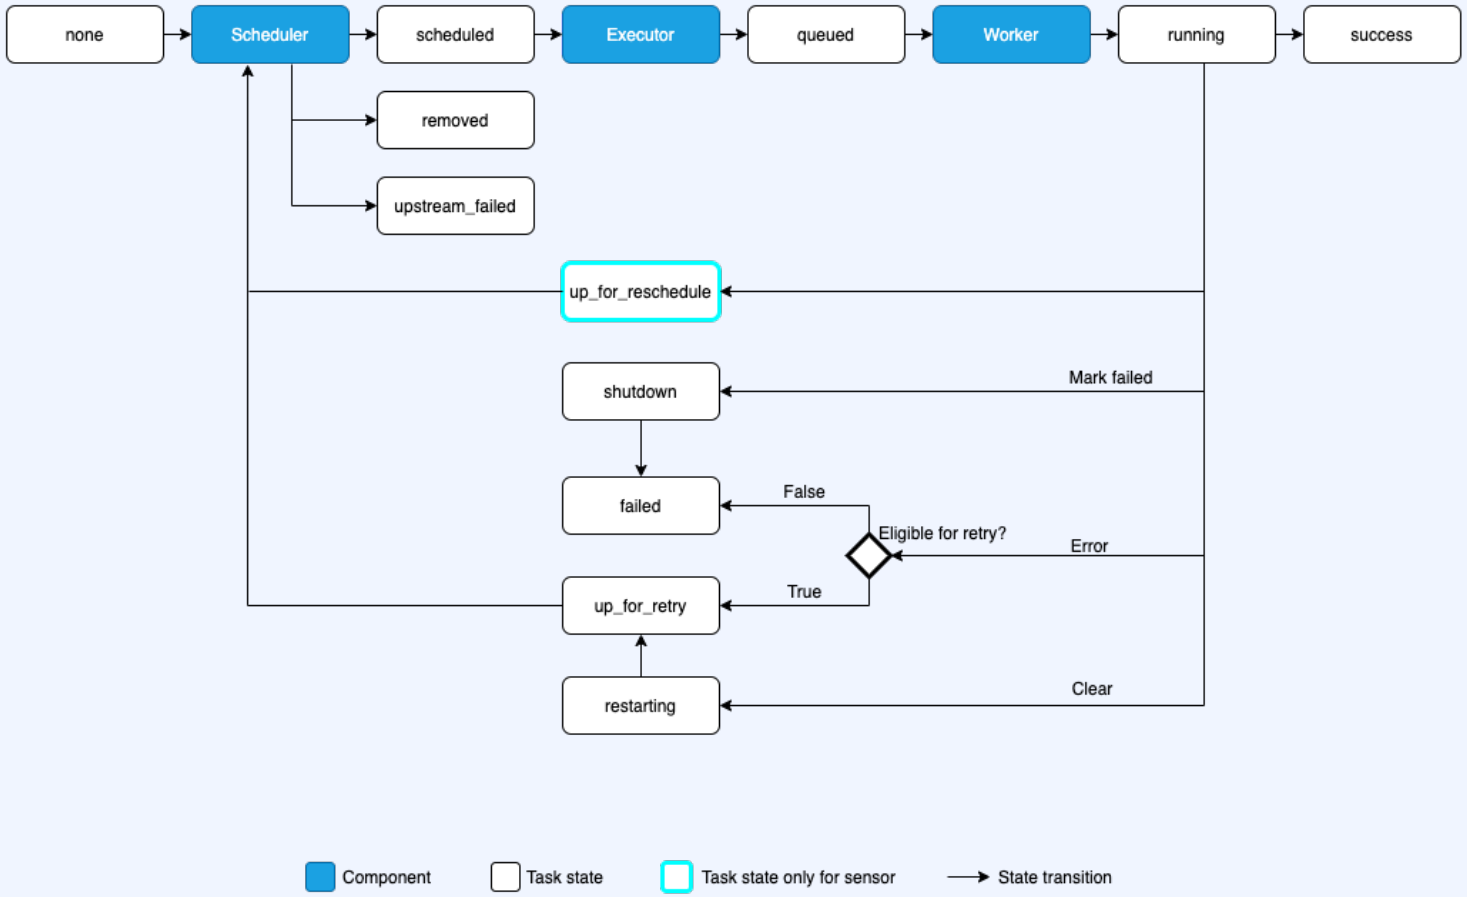

01. Components

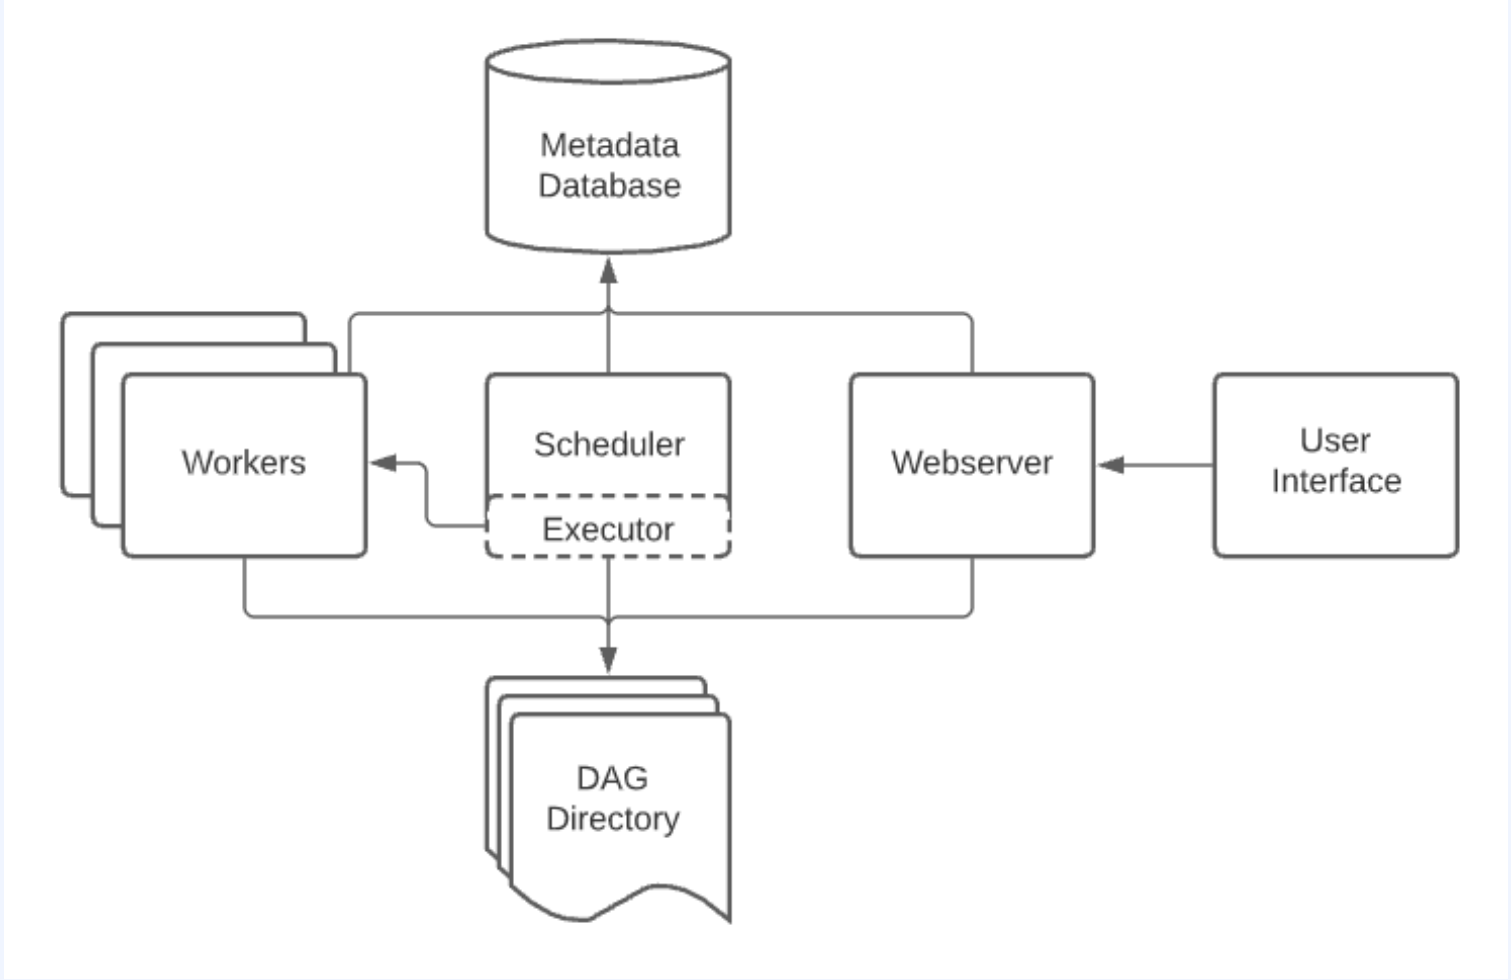

01. Workloads

- Task의 모든 type은 내부적으로 BaseOperator의 subclass
- Three common types
  - Operators
  - Sensors
  - TaskFlow-decorated @task

01. TaskFlow 예제

```
@task
def get_ip():
    return my_ip_service.get_main_ip()


@task
def compose_email(external_ip):
    return {
        'subject':f'Server connected from {external_ip}',
        'body': f'Your server executing Airflow is connected from the external IP
{external_ip}<br>'
    }
```

01. TaskFlow 예제

```
email_info = compose_email(get_ip())

EmailOperator(
    task_id='send_email',
    to='example@example.com',
    subject=email_info['subject'],
    html_content=email_info['body']
)
```

01. Control Flow

- DAGs
  - run many times, and multiple runs of them can happen in parallel
  - always including an interval they are “running for” (the data interval)
- Tasks
  - have dependencies declared on each other
    - first_task >> second_task

01. Circular Dependency

```
from airflow.operators.dummy_operator import DummyOperator
from datetime import datetime

with DAG('circular_dependency_dag', start_date=datetime(2022, 1, 1)) as dag:
    task1 = DummyOperator(task_id='task1')
    task2 = DummyOperator(task_id='task2')

    task1 >> task2 >> task1
```

01. Circular Dependency

```
with DAG('DAG1', default_args=default_args, schedule_interval=None) as dag:
    trigger_dag2 = TriggerDagRun(
                task_id='trigger_dag2’,
                trigger_dag_id="DAG2",
                conf={"message": "Hello from DAG1"})

with DAG('DAG2', default_args=default_args, schedule_interval=None) as dag:
    trigger_dag1 = TriggerDagRun(
                task_id='trigger_dag1’,
                trigger_dag_id="DAG1",
                conf={"message": "Hello from DAG2"} )
```

01. Pass data between tasks

- Xcoms
- Storage Service

01. Communication with external services

- Connections & Hooks
- Pools

01. User Interface

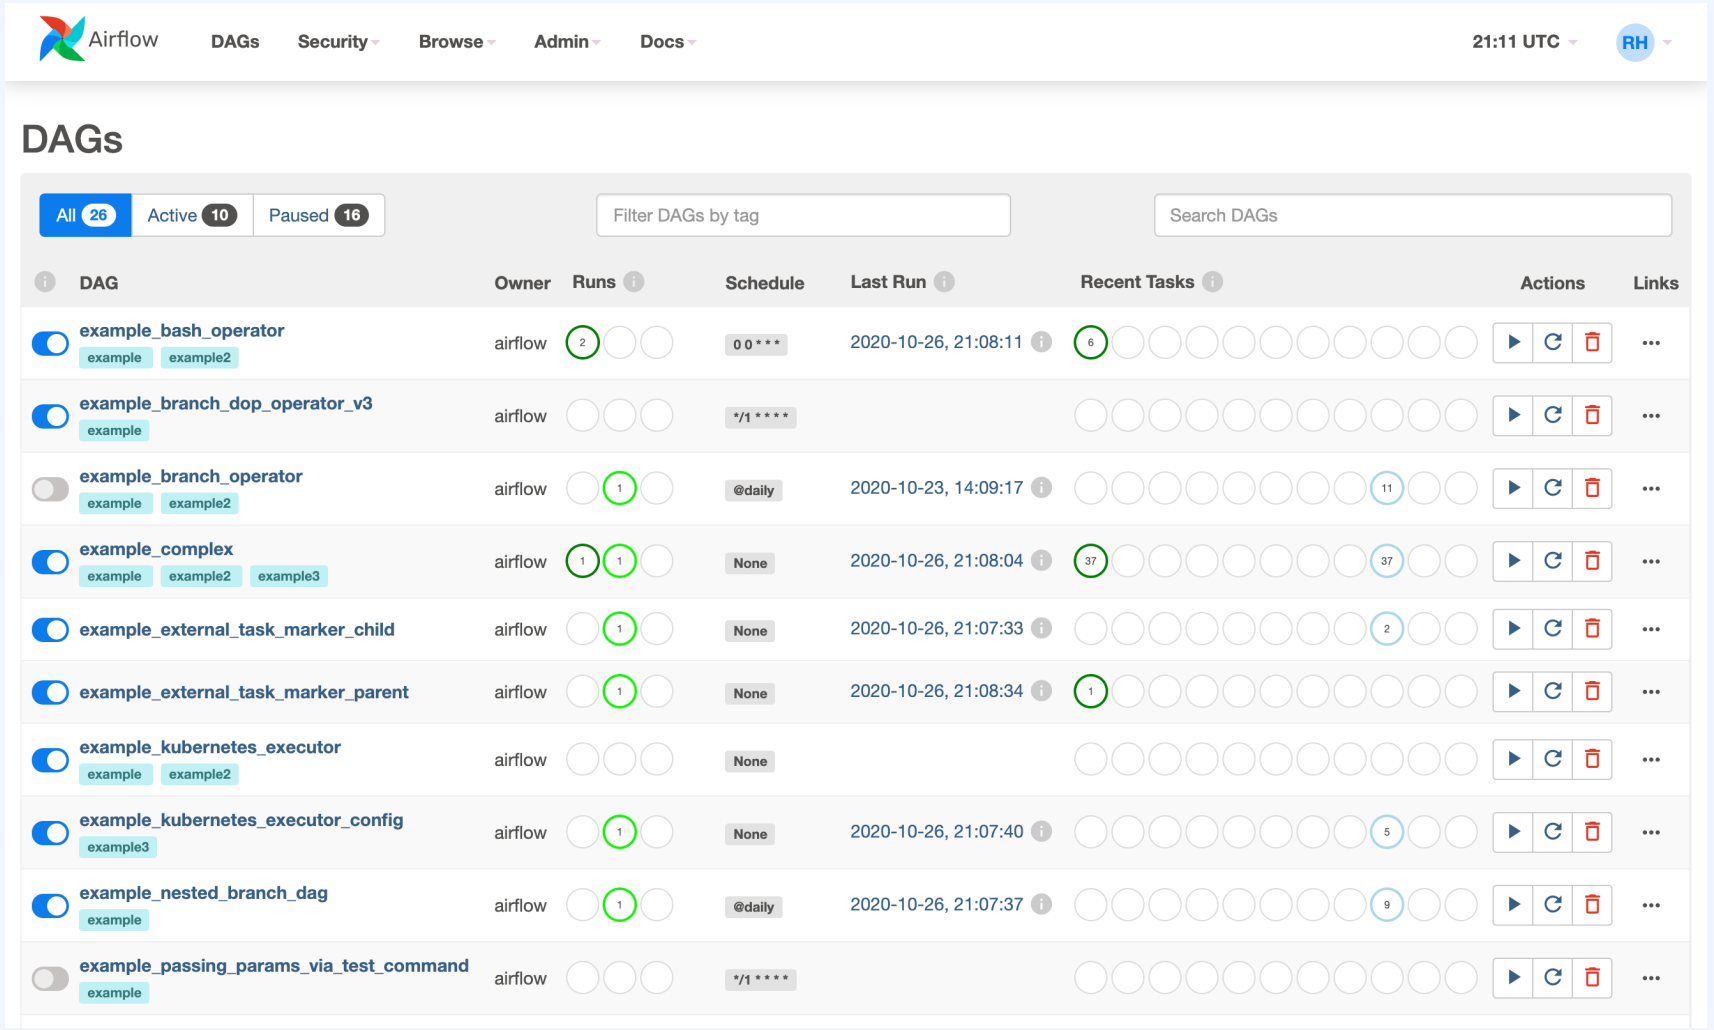

Chapter 2.

01. DAG

01. DAG

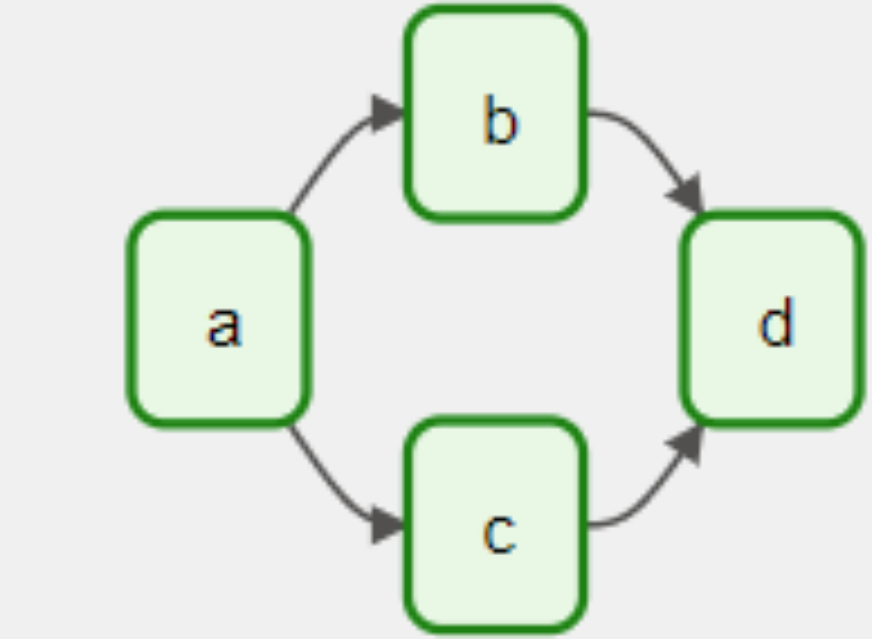

01. Declaring a DAG – Context Manager

```
with DAG(
    dag_id="my_dag_name",
    start_date=datetime.datetime(2021, 1, 1),
    schedule="@daily",
):
    EmptyOperator(task_id="task")
```

01. Declaring a DAG – Context Manager

```
my_dag = DAG(
    dag_id="my_dag_name",
    start_date=datetime.datetime(2021, 1, 1),
    schedule="@daily",
)
EmptyOperator(task_id="task", dag=my_dag)
```

01. Declaring a DAG – Context Manager

```
@dag(start_date=datetime.datetime(2021, 1, 1), schedule="@daily")
def generate_dag():
    EmptyOperator(task_id="task")

generate_dag()
```

01. Loading DAGs

```
dag_1 = DAG('this_dag_will_be_discovered')

def my_function():
    dag_2 = DAG('but_this_dag_will_not')

my_function()
```

01. Running DAGs

```
with DAG("my_daily_dag", schedule="@daily"):
    ...
```

01. DAG Assignment

```
with DAG("my_dag"):
    some_operator = SomeOperator(task_id="some_task")
```

01. Edge Labels

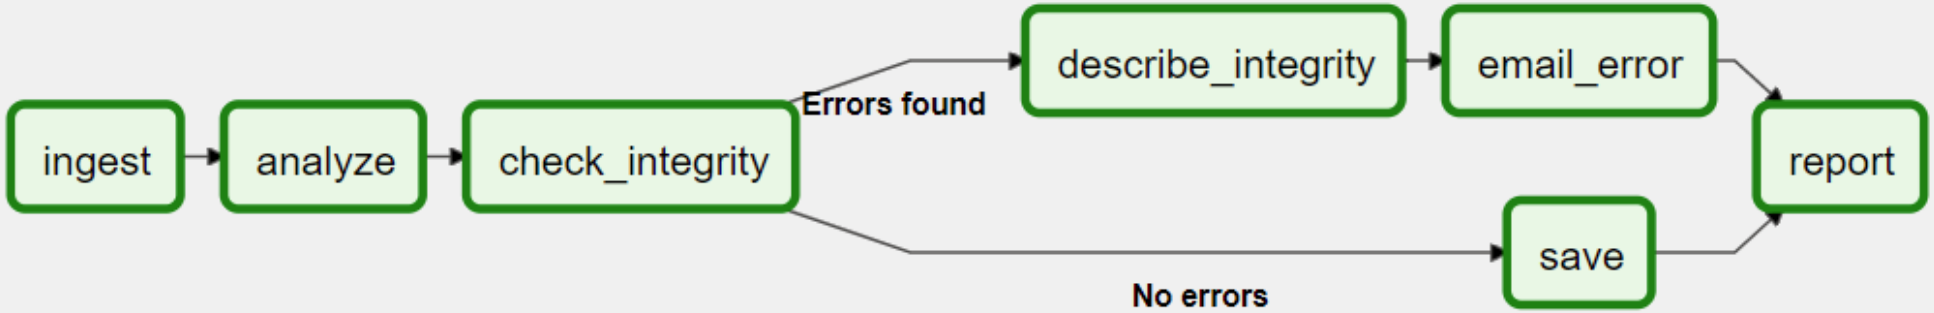

01. Edge Labels

```
ingest = EmptyOperator(task_id="ingest")
analyse = EmptyOperator(task_id="analyze")
check = EmptyOperator(task_id="check_integrity")
describe = EmptyOperator(task_id="describe_integrity")
error = EmptyOperator(task_id="email_error")
save = EmptyOperator(task_id="save")
report = EmptyOperator(task_id="report")

ingest >> analyse >> check
check >> Label("No errors") >> save >> report
check >> Label("Errors found") >> describe >> error >> report
```

Chapter 2.

02. Tasks

02. Tasks

- Operators
- Sensors
- Taskflow – decorated @task

02. Relations

>> or <<

- first_task >> second_task >> [third_task, fourth_task]

upstraeam() or downstream()

- first_task.set_downstream(second_task)
- third_task.set_upstream(second_task)

02. Task Instances

- none
- scheduled
- queued
- running
- success
- shutdown
- restarting

02. Task Instances

- failed
- skipped
- upstream_failed
- up_for_retry
- up_for_reschedule
- deferred
- removed

02. Task Instances

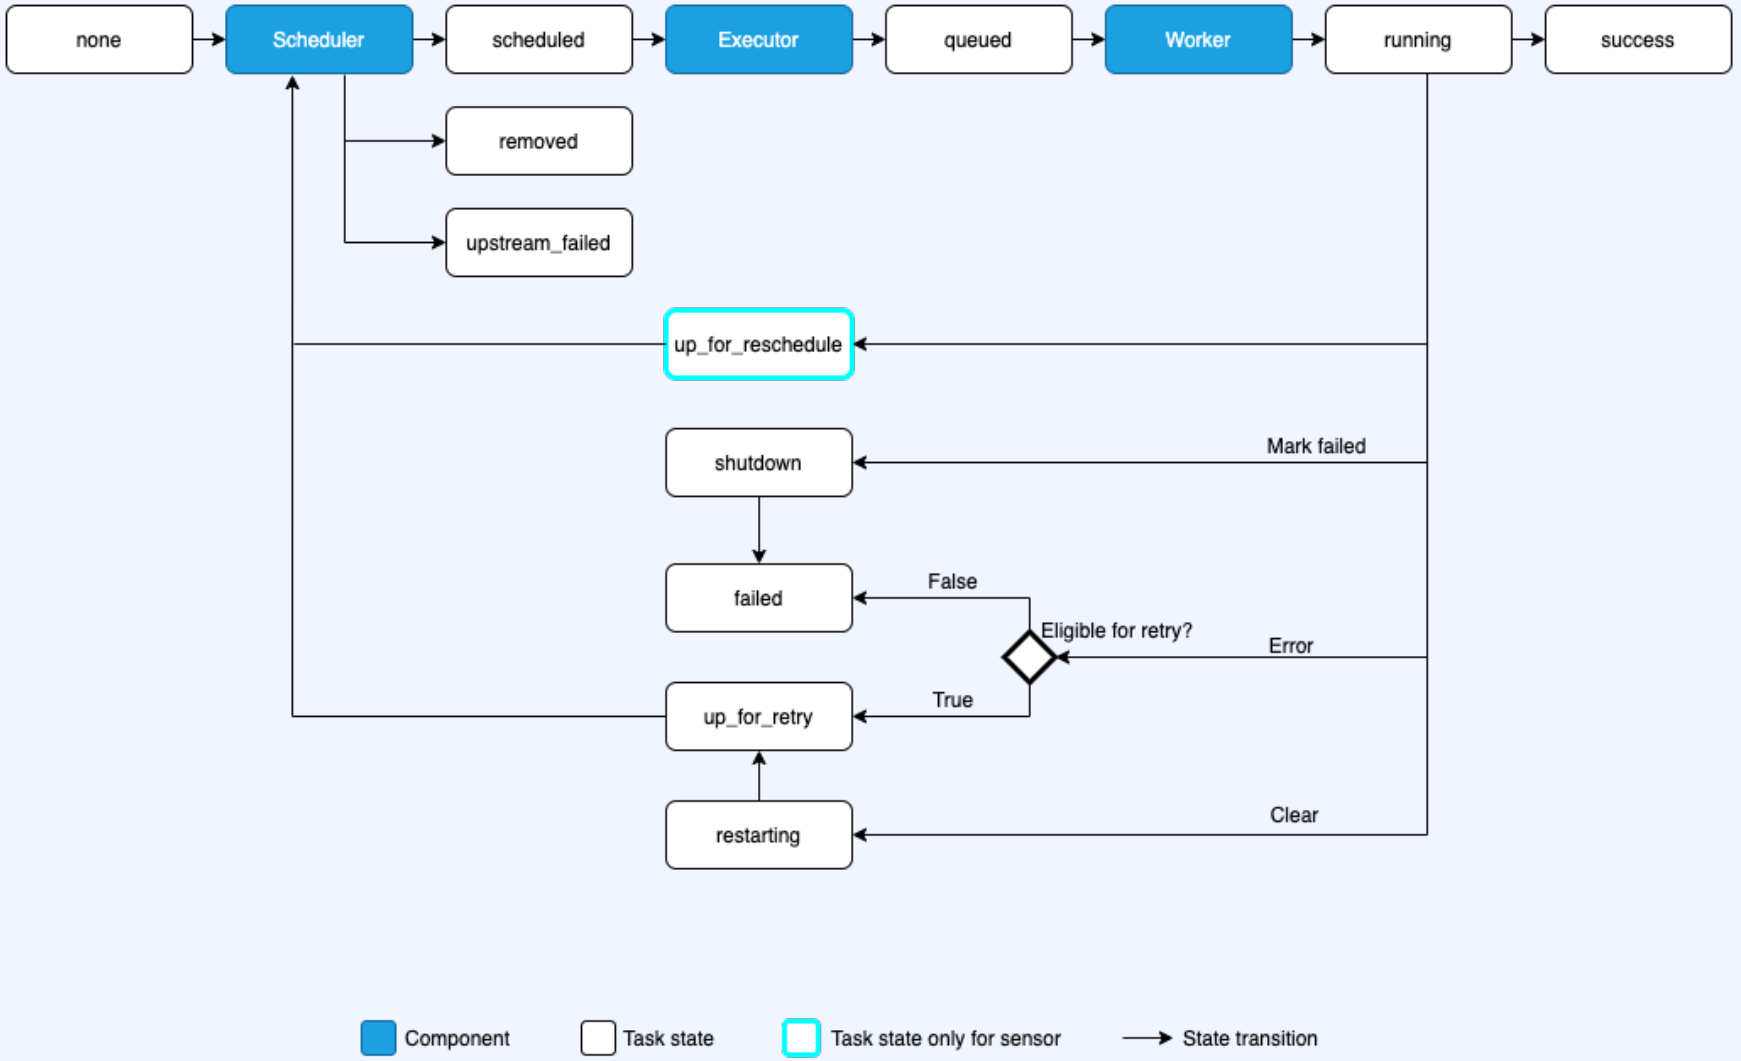

02. Special Exceptions

```
def skip_task():
    raise AirflowSkipException("This task is skipped!")
def fail_task():
    raise AirflowFailException("This task has failed!")

skip_task_operator = PythonOperator( task_id='skip_task',
python_callable=skip_task, dag=dag, )
fail_task_operator = PythonOperator( task_id='fail_task',
python_callable=fail_task, dag=dag, )
```

02. Zombie/Undead Tasks

- Zombie tasks
- Undead tasks

02. Executor Configuration

```
MyOperator(...,
    executor_config={
        "KubernetesExecutor":
            {"image": "myCustomDockerImage"}
    }
)
```

02. DAG & Task Documentation

```
""" ### My great DAG """

dag = DAG( "my_dag", start_date=pendulum.datetime(2021, 1, 1, tz="UTC"),
schedule="@daily", catchup=False, )
dag.doc_md = __doc__

t = BashOperator("foo", dag=dag)
t.doc_md = """
  \ #Title" Here's a [url](www.airbnb.com)
"""
```

02. DAG & Task Documentation

- doc
- doc_json
- doc_yaml
- doc_md
- doc_rst

Chapter 2.

03. Control Flow

03. Branching

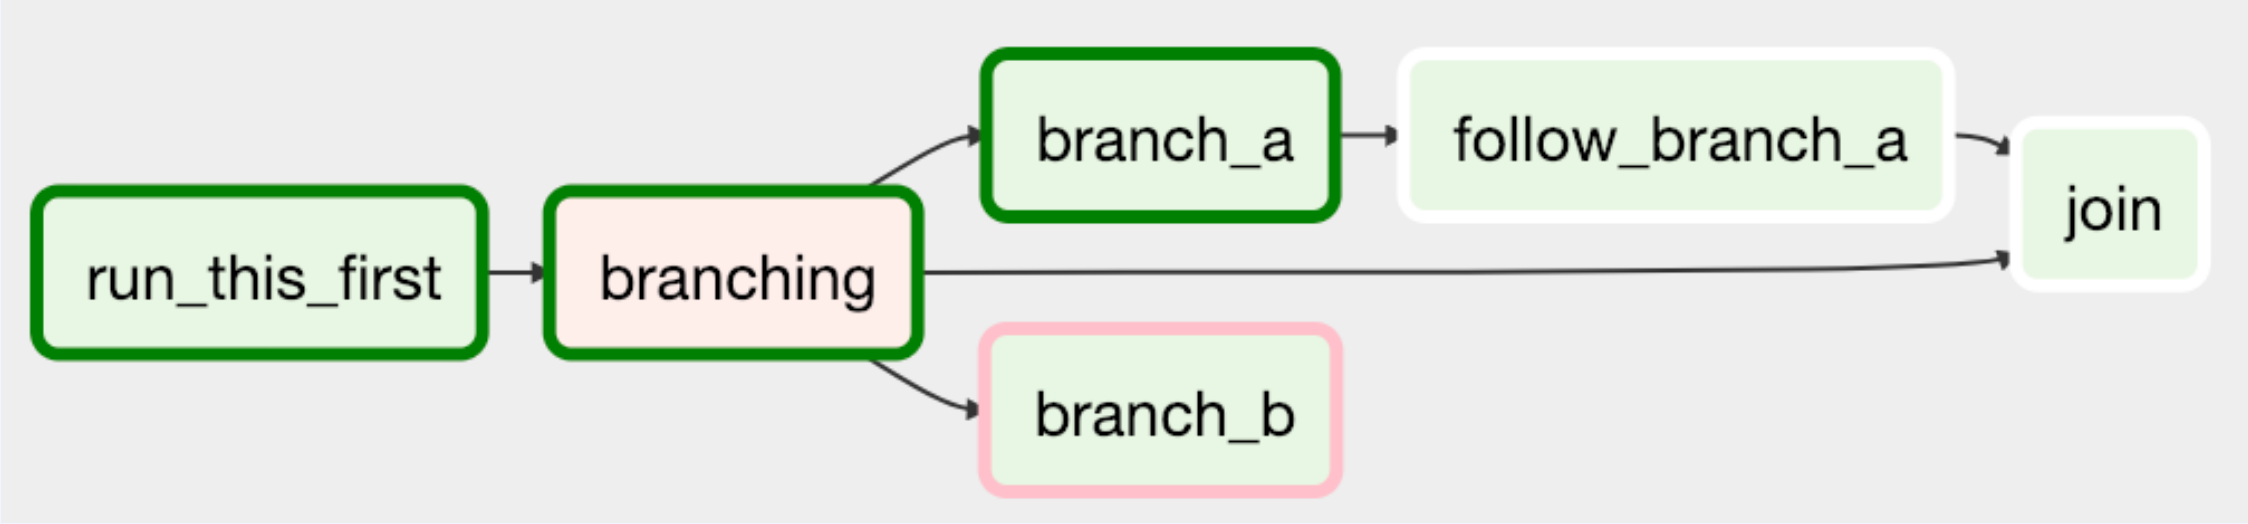

03. Branching

```
@task.branch(task_id="branch_task")
def branch_func(ti=None):
    xcom_value = int(ti.xcom_pull(task_ids="start_task"))
    if xcom_value >= 5:
            return "continue_task"
    elif xcom_value >= 3:
            return "stop_task"
    else: return None

branch_op = branch_func()
start_op >> branch_op >> [continue_op, stop_op]
```

03. Branching

```
class MyBranchOperator(BaseBranchOperator):
    def choose_branch(self, context):
        if context['data_interval_start'].day == 1:
            return ['daily_task_id', 'monthly_task_id']
        elif context['data_interval_start'].day == 2:
            return 'daily_task_id'
        else:
            return None
```

03. Latest Only

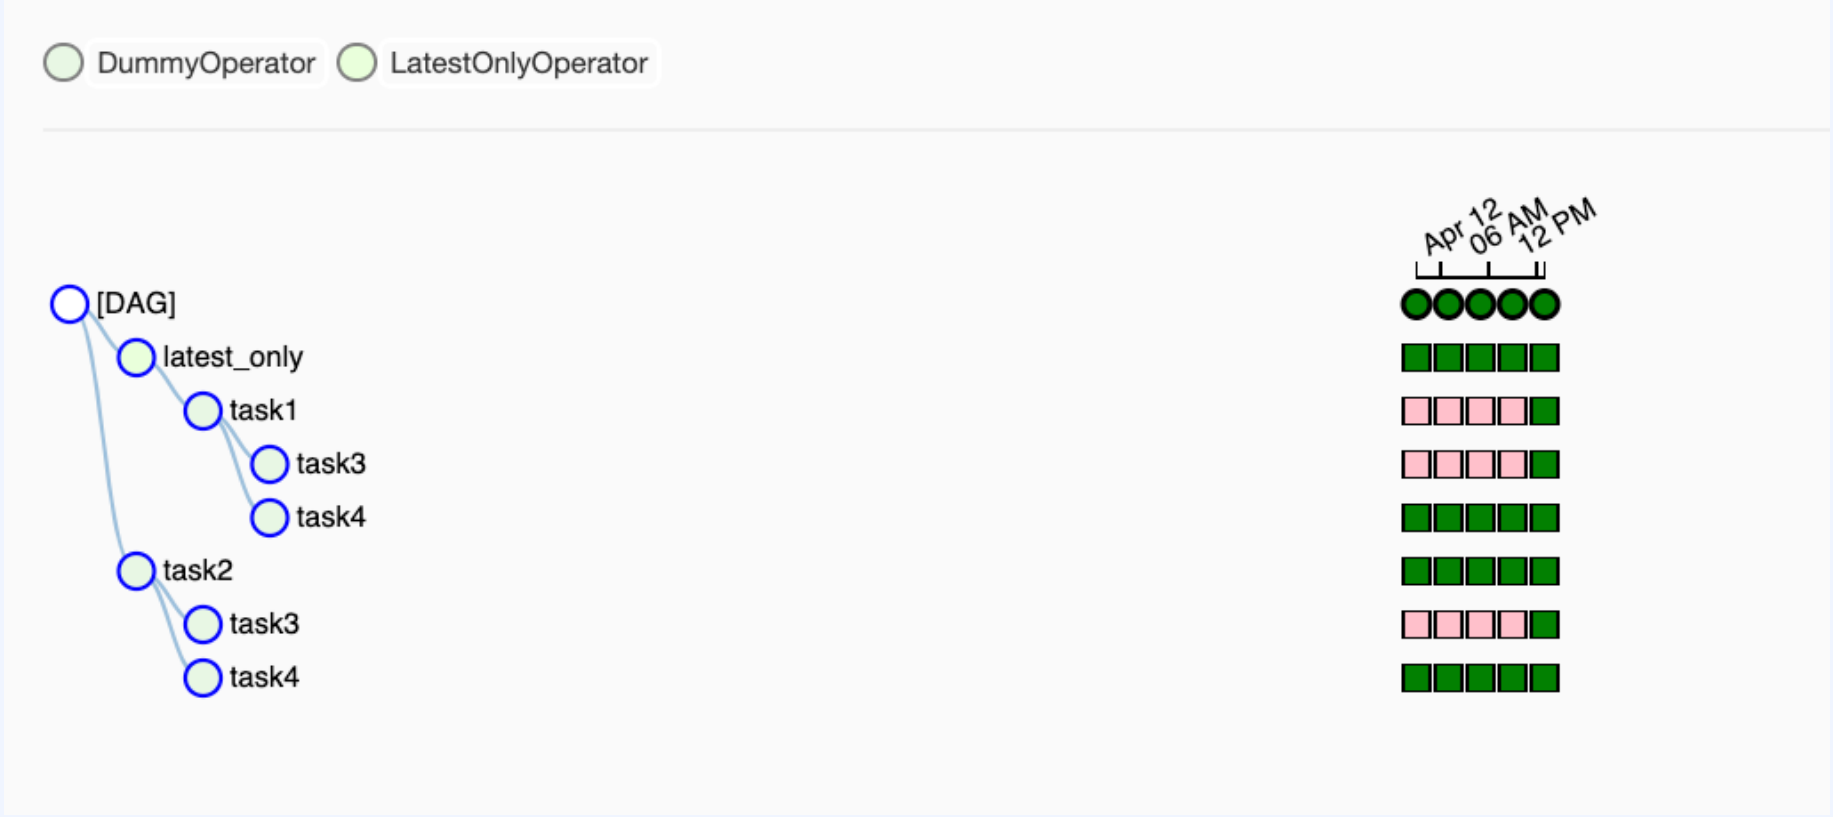

03. Branching

```
latest_only = LatestOnlyOperator(task_id="latest_only")
task1 = EmptyOperator(task_id="task1")
task2 = EmptyOperator(task_id="task2")
task3 = EmptyOperator(task_id="task3")
task4 = EmptyOperator(task_id="task4", trigger_rule=TriggerRule.ALL_DONE)

latest_only >> task1 >> [task3, task4]
task2 >> [task3, task4]
```

03. Depends On Past

- ‘depends_on_past’: True

03. Trigger Rules

- all_success (기본값)
- all_failed
- all_done
- all_skipped
- one_failed
- one_success
- one_done
- none_failed
- none_failed_min_one_success
- none_skipped
- always

03. Trigger Rules

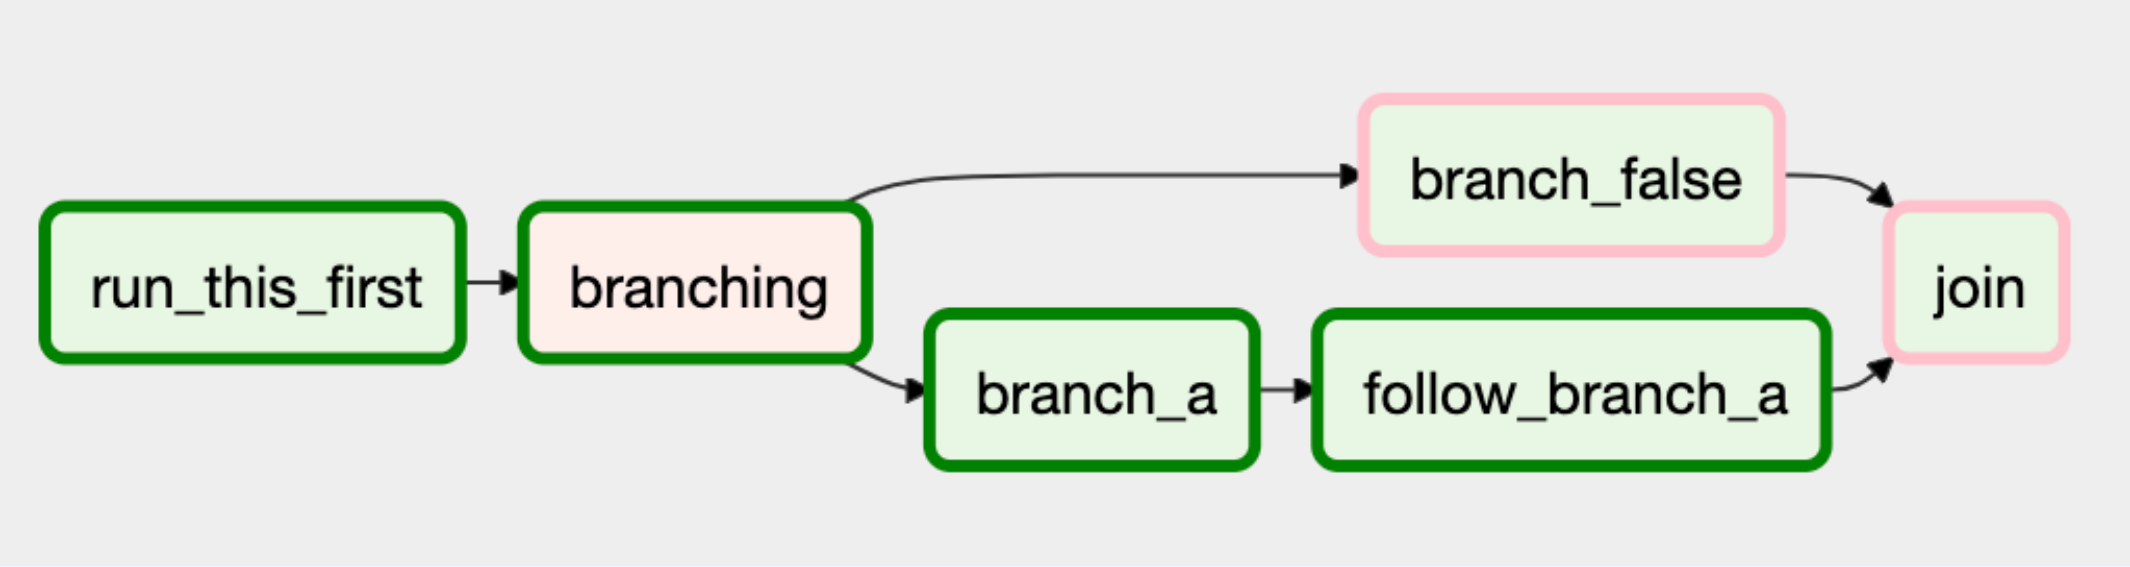

03. Trigger Rules

```
join = EmptyOperator(task_id="join",
    trigger_rule=TriggerRule.NONE_FAILED_MIN_ONE_SUCCESS,
    dag=dag)

run_this_first >> branching
branching >> branch_a >> follow_branch_a >> join
branching >> branch_false >> join
```

03. Trigger Rules

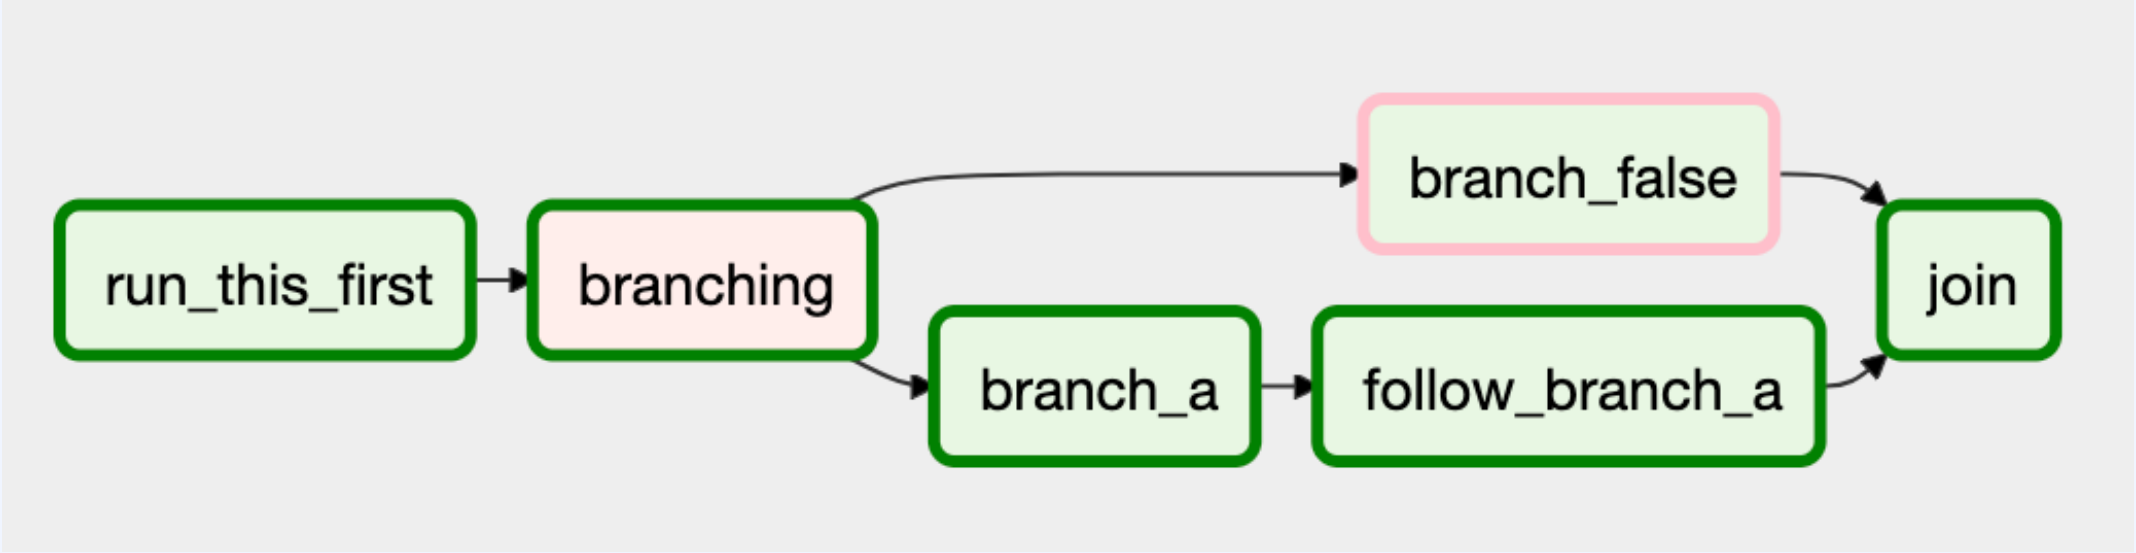

Chapter 2.

04. Timetables

04. DAG Runs

- An instantiation of the DAG **in time**
- The status of the DAG Run depends on the tasks states

04. DAG Run Status

- **success** if all of the leaf nodes states are either success or skipped
- **failed** if any of the leaf nodes state is either failed or upstream_failed

04. Cron Presets

- None
- @once
- @continuously
- @hourly: 0 * * * *
- @daily: 0 0 * * *
- @weekly: 0 0 * * 0
- @monthly: 0 0 1 * *
- @quarterly: 0 0 1 */3 *
- @yearly: 0 0 1 1 *

04. Data Interval

- execution_date(-> logical date)
- data_interval_start
- data_intercal_end

04. Data Interval

- 매일 실행되는 DAG의 경우:
  - logical_date (execution_date): 2021-01-02 00:00:00
  - data_interval_start: 2021-01-02 00:00:00
  - data_interval_end: 2021-01-03 00:00:00

- 매일 실행되는 DAG가 2일 전의 데이터를 처리
  - logical_date (execution_date): 2021-01-03 00:00:00
  - data_interval_start: 2021-01-01 00:00:00
  - data_interval_end: 2021-01-02 00:00:00

04. Timetables

- DAGs are driven by its internal “timetable”
- The timetable determines the data interval and the logical date

04. Custom Timetables

- Data intervals with “holes” between.
- Run tasks at different times each day.
- Schedules not following the Gregorian calendar.
- Rolling windows, or overlapping data intervals.

04. Implement Custom Timetables

```
class PassWeekendTimetable(Timetable):
    def next_dagrun_info( self, *, last_automated_dagrun: Optional[datetime],
**kwargs ) -> Optional[DagRunInfo]:
            if not last_automated_dagrun:
                    return self._first_dagrun()
            next_start = last_automated_dagrun + timedelta(days=1)
            if next_start.weekday() >= 5:
                    next_start += timedelta(days=2)
                    return DataInterval(
                            start=next_start, end=next_start + timedelta(days=1), )
```

04. Built-in Timetables – Events Timetable

```
@dag(
 schedule=EventsTimetable(
  event_dates=[
   pendulum.datetime(2022, 4, 5, 8, 27, tz="America/Chicago"),
   pendulum.datetime(2022, 4, 17, 8, 27, tz="America/Chicago"),
   pendulum.datetime(2022, 4, 22, 20, 50, tz="America/Chicago"),
  ],
  description="My Team's Baseball Games",
  restrict_to_events=False, ), ..., ) def example_dag(): pass
```

04. Differences between the two cron timetables

- CronTriggerTimetable: does not care the idea of data interval
- CronDataIntervalTimetable: does care the idea of data interval

- If catchup is False
  - CronTriggerTimetable: a new DAG run after the current time
  - CronDataIntervalTimetable: a new DAG run before the current time

Chapter 2.

05. Catchup & Backfill

05. Re-run DAG

- Re-run DAG != Create a new DAG run

05. Catchup

- start_date, end_date, none-dataset schedule
  -> defines a series of intervals: the scheduler turns into individual DAG runs
- scheduler kick off a DAG Run for any data interval
  that has not been run since the last data interval

- In code, catch_up=False
- In config, catch_up_default=False

05. Catchup

```
dag = DAG(
    "tutorial",
    default_args={ "depends_on_past": True,
                "retries": 1,
                "retry_delay": datetime.timedelta(minutes=3), },
    start_date=pendulum.datetime(2015, 12, 1, tz="UTC"),
    description="A simple tutorial DAG",
    schedule="@daily",
    catchup=False, )
```

05. Backfill

```
airflow dags backfill \
    --start-date START_DATE \
    --end-date END_DATE \
    dag_id
```

05. To re-run Tasks

- Re-run the tasks by clearing them for the scheduled date
- Clearing a task instance doesn't delete the task instance record.
  - Instead, it updates **max_tries to 0**
  - sets the current **task instance state to None**, which causes the task to re-run.
- In the Tree or Graph views, click **Clear**

05. Re-run Options

- Past
- Future
- Upstream
- Downstream
- Recursive
- Failed

05. Re-run CLI

```
airflow tasks clear dag_id \
    --task-regex task_regex \
    --start-date START_DATE \
    --end-date END_DATE
```

Chapter 2.

06. Timeout

06. execution_time

- execution_timeout attribute to a **datetime.timedelta** value
- execution_timeout is breached,

  the task times out and **AirflowTaskTimeout** is raised.

06. timeout (Sensor / Reschedule mode)

- timeout controls the maximum time allowed for the sensor to succeed
- If timeout is breached, **AirflowSensorTimeout** will be raised

06. SFTPSensor 예제

```
sensor = SFTPSensor(
    task_id="sensor",
    path="/root/test",
    execution_timeout=timedelta(seconds=60),
    timeout=3600,
    retries=2,
    mode="reschedule",
)
```

06. SLAs

- Send the alarm when running the task

06. Set an SLA

```
def sla_miss_callback(context):
    pass


task_with_sla = DummyOperator(
    task_id='task_with_sla',
    sla=timedelta(hours=2),.
    sla_miss_callback=sla_miss_callback,
    dag=dag
)
```

06. Disable SLA

[core]

check_slas = False

Chapter 2.

07. Callbacks & Notifier

07. Callbacks

- 중요성
  - 상태 모니터링
  - 자동화

- 주의사항
  - Worker 실행에 의한 상태 변화만 적용
  - Callback의 오류 Logging
    - $AIRFLOW_HOME/logs/scheduler/latest/PROJECT/DAG_FILE.py.log

07. Callback Types

- on_success_callback
- on_failure_callback
- sla_miss_callback
- on_retry_callback
- on_execute_callback

07. Callback 예제

```
with DAG( dag_id="example_callback", schedule=None,
    start_date=pendulum.datetime(2021, 1, 1, tz="UTC"),
    dagrun_timeout=datetime.timedelta(minutes=60), catchup=False,
    on_success_callback=None,
    on_failure_callback=task_failure_alert,
    tags=["example"], ):
task1 = EmptyOperator(task_id="task1")
task2 = EmptyOperator(task_id="task2")
task3 = EmptyOperator(task_id="task3",
    on_success_callback=[dag_success_alert])
task1 >> task2 >> task3
```

07. BaseNotifier

- An abstract class
- Sending notifications using the various **on_*_callback**

07. Extend the BaseNotifier

```
class MyNotifier(BaseNotifier):
    template_fields = ("message",)

    def __init__(self, message):
        self.message = message

    def notify(self, context):
        # Send notification here, below is an example
        title = f"Task {context['task_instance'].task_id} failed"
        send_message(title, self.message)
```

07. Using a notifier

```
with DAG(dag_id="example_notifier",
    start_date=datetime(2022, 1, 1),
    schedule_interval=None,
    on_success_callback=MyNotifier(message="Success!"),
    on_failure_callback=MyNotifier(message="Failure!"), ):

    task = BashOperator(
        task_id="example_task",
        bash_command="exit 1",
        on_success_callback=MyNotifier(message="Task Succeeded!"),
    )
```

07. Notifications

- Amazon: ChimeNotifier
- Apprise: AppriseNotifier
- Discord: DiscordNotifier
- Pagerduty: PagerdutyNotifier
- Slack: SlackNotifier
- Simple Mail Transfer Protocol (SMTP): SmtpNotifier

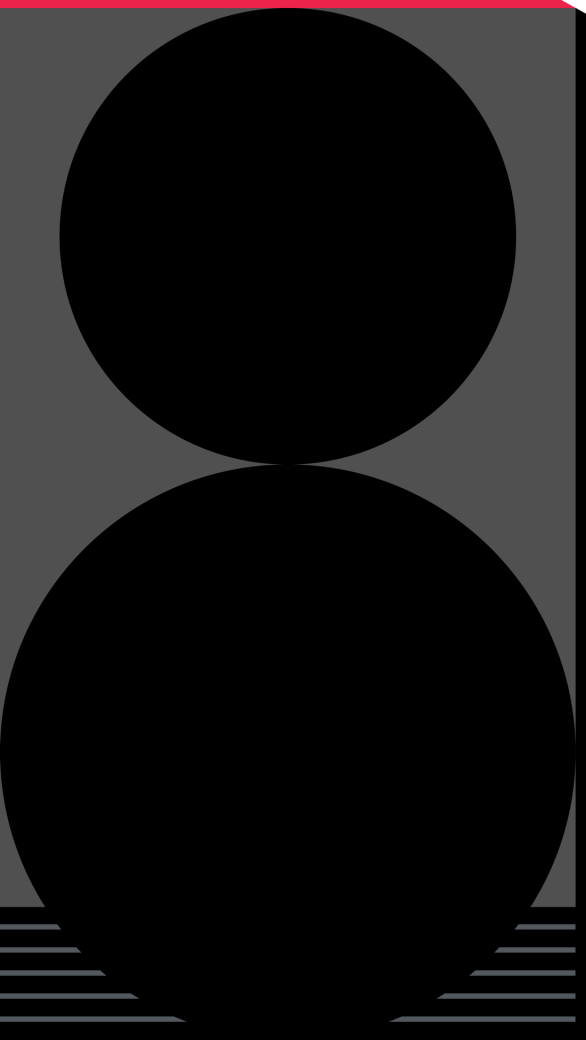

Chapter 2.

08. Pools

08. Pools

- **limit the execution parallelism** on arbitrary sets of tasks
- managed in the UI (Menu -> Admin -> Pools)
  - a name and assigning it a number of worker slots

08. pool parameter

```
aggregate_db_message_job = BashOperator(
    task_id="aggregate_db_message_job",
    execution_timeout=timedelta(hours=3),
    pool="ep_data_pipeline_db_msg_agg",
    bash_command=aggregate_db_message_job_cmd,
    dag=dag,
)
aggregate_db_message_job.set_upstream(wait_for_empty_queue)
```

08. Pools

- Number of slots occupied by a task == pool_slots
- As slots free up, queued tasks start running based on the **Priority Weights**
- If tasks are not given a pool -> default_pool(128 slots)

08. Priority Weights

- **priority_weight**(default: 1)
  - defines priorities in the executor queue
  - calculated based on its **weight_rule**

08. Weighting Methods

- airflow.utils.WeightRule
  - downstream (default)
  - upstream
  - absolute

08. Weighting Methods - Example

```
task_1 = EmptyOperator( task_id='task_1’,
    weight_rule=WeightRule.DOWNSTREAM, )
task_2 = EmptyOperator( task_id='task_2’,
    weight_rule=WeightRule.UPSTREAM, )
task_3 = EmptyOperator( task_id='task_3’,
    weight_rule=WeightRule.ABSOLUTE,
    priority_weight=10, )

task_1 >> task_2 >> task_3
```

08. Weighting Methods

- a task -> a single pool slot (default)
- pool_slots: can be configured to occupy more

08. Weighting Methods - Example

```
BashOperator( task_id="heavy_task",
    bash_command="bash backup_data.sh",
    pool_slots=2, pool="maintenance", )
BashOperator( task_id="light_task1",
    bash_command="bash check_files.sh",
    pool_slots=1, pool="maintenance", )
BashOperator( task_id="light_task2",
    bash_command="bash remove_files.sh",
    pool_slots=1, pool="maintenance", )
```

Chapter 2.

09. Connections & Hook

09. Connections

- A set of parameters(username, password and hostname)
  with the type of system
  and a unique name(conn_id)
- Managed via the UI or the CLI
- Customizable connection storage and backend options

09. Managing Connections

- In environment variables
  - AIRFLOW_CONN_{CONN_ID}
- In an external Secrets Backend
  - AWS Secrets Manager, HashiCorp Vault, Google Cloud Secret Manager
- In the Airflow metadata database (using the CLI or web UI)

09. Using Connection

```
conn = BaseHook.get_connection('<conn_id>')
print(conn.host)
```

09. Using Connection

```
pg_hook = PostgresHook('<conn_id>')

records = pg_hook.get_records('SELECT * FROM your_table')
print(records)
```

09. Using Connection

```
print_host = BashOperator(
    task_id='print_host', bash_command='echo {{ conn.<conn_id>.host }}'
)
```

09. Hook

- A **high-level interface** to an external platform
- **Building blocks** that Operators are built out of
- Have a **default conn_id**
  - In PostgresHook, default conn_id is postgres_default

Chapter 2.

10. SubDAGs & TaskGroups

10. SubDAGs

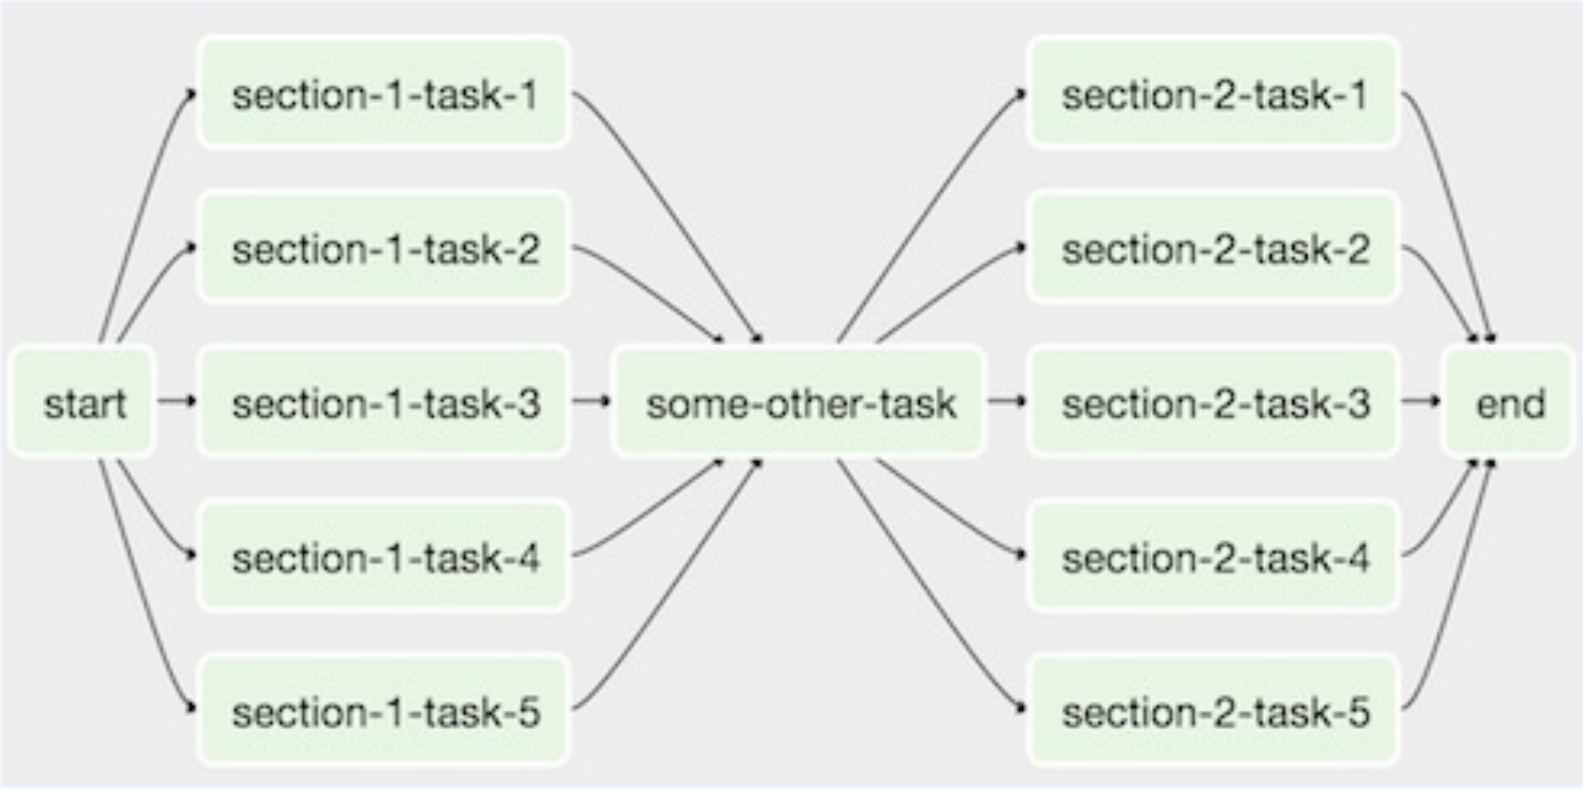

10. SubDAGs

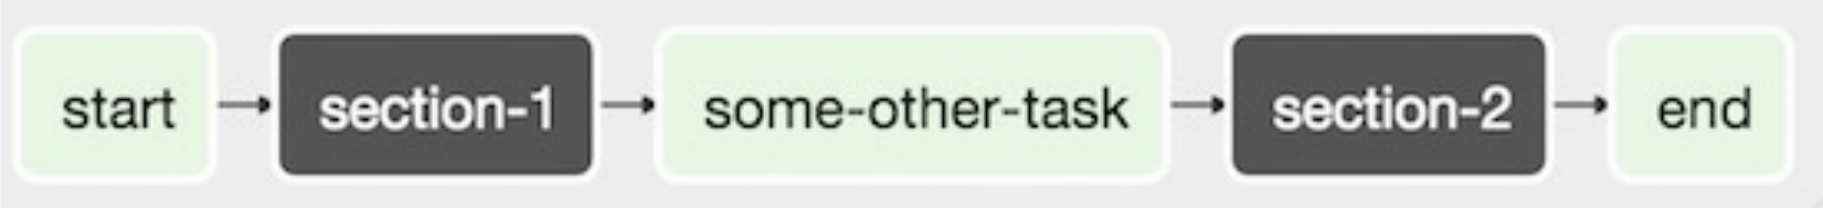

10. subdag() factory method

```
def subdag(parent_dag_name, child_dag_name, args) -> DAG:
    dag_subdag = DAG(
            dag_id=f"{parent_dag_name}.{child_dag_name} ”,
            ··· )
    for i in range(5):
            EmptyOperator(···, dag=dag_subdag, )
    return dag_subdag

subdag_task = SubDagOperator( task_id='subdag_task’,
    subdag=subdag('main_dag', 'subdag_task', args),
    dag=main_dag, )
```

10. TaskGroups

- Organize tasks into hierarchical groups in Graph
- **Purely a UI grouping concept**
  - https://airflow.apache.org/docs/apache-airflow/stable/_images/task_group.gif

10. TaskGroups Example

```
from airflow.decorators import task_group


@task_group()
def group1():
    task1 = EmptyOperator(task_id="task1")
    task2 = EmptyOperator(task_id="task2")

task3 = EmptyOperator(task_id="task3")

group1() >> task3
```

10. TaskGroups Example – default_args

```
with DAG( dag_id="dag1", start_date=datetime.datetime(2016, 1, 1),
    schedule="@daily", default_args={"retries": 1}, ):
    @task_group(default_args={"retries": 3})
    def group1():
            """This docstring will become the tooltip for the TaskGroup."""
            task1 = EmptyOperator(task_id="task1")
            task2 = BashOperator(task_id="task2", bash_command="echo Hello
World!", retries=2)
            print(task1.retries) # 3
            print(task2.retries) # 2
```

10. TaskGroups vs SubDAGs

- SubDAG is deprecated hence TaskGroup is **always the preferred choice**.

Chapter 2.

11. XComs & Variables

11. XComs

- Short for ‘cross-communication’
- Let Tasks talk to each other

11. XComs

- Identifier
  - key(name), task_id, dag_id
- Value
  - serializable
  - only designed for small amounts of data
- Push & Pull
  - xcom_push, xcom_pull
  - If do_xcom_push == True (default)
    - retrun_value: auto-push result into XCom key
      - @task

11. XComs

- Templates
  - SELECT * FROM {{ task_instance.xcom_pull(task_ids='foo', key='table_name') }}
- vs Variables
  - XComs: per-task-instance, designed for communication
  - Variables: global, designed for overall configuration and value sharing

11. Custom Xcoms Backends

- Airflow metadate db (default)

11. Custom Xcoms Backends

- Interchangeable
  - xcom_backend
- Implement
  - BaseXCom: serialize_value, deserialize_value
- Rendered for UI
  - orm_deserialized_value
- Lifecycle
  - clear

11. Custom Xcom Backends - Container

```
from airflow.models.xcom import XCom
print(XCom.__name__)

from airflow.settings import conf
conf.get("core", "xcom_backend")
```

11. Variables

- **Runtime** configuration concept
  - A general key/value store **that is global**
  - **Can be quried** from tasks
- Set via Airflow's user interface or bulk uploaded as a JSON file

11. Use Variables – get()

```
from airflow.models import Variable

# Normal call style
foo = Variable.get("foo")

# Auto-deserializes a JSON value
bar = Variable.get("bar", deserialize_json=True)

# Returns the value of default_var (None) if the variable is not set
baz = Variable.get("baz", default_var=None)
```

11. Use Variables – templates

```
# Raw value
echo {{ var.value.<variable_name> }}

# Auto-deserialize JSON value
echo {{ var.json.<variable_name> }}
```

11. Managing Variables

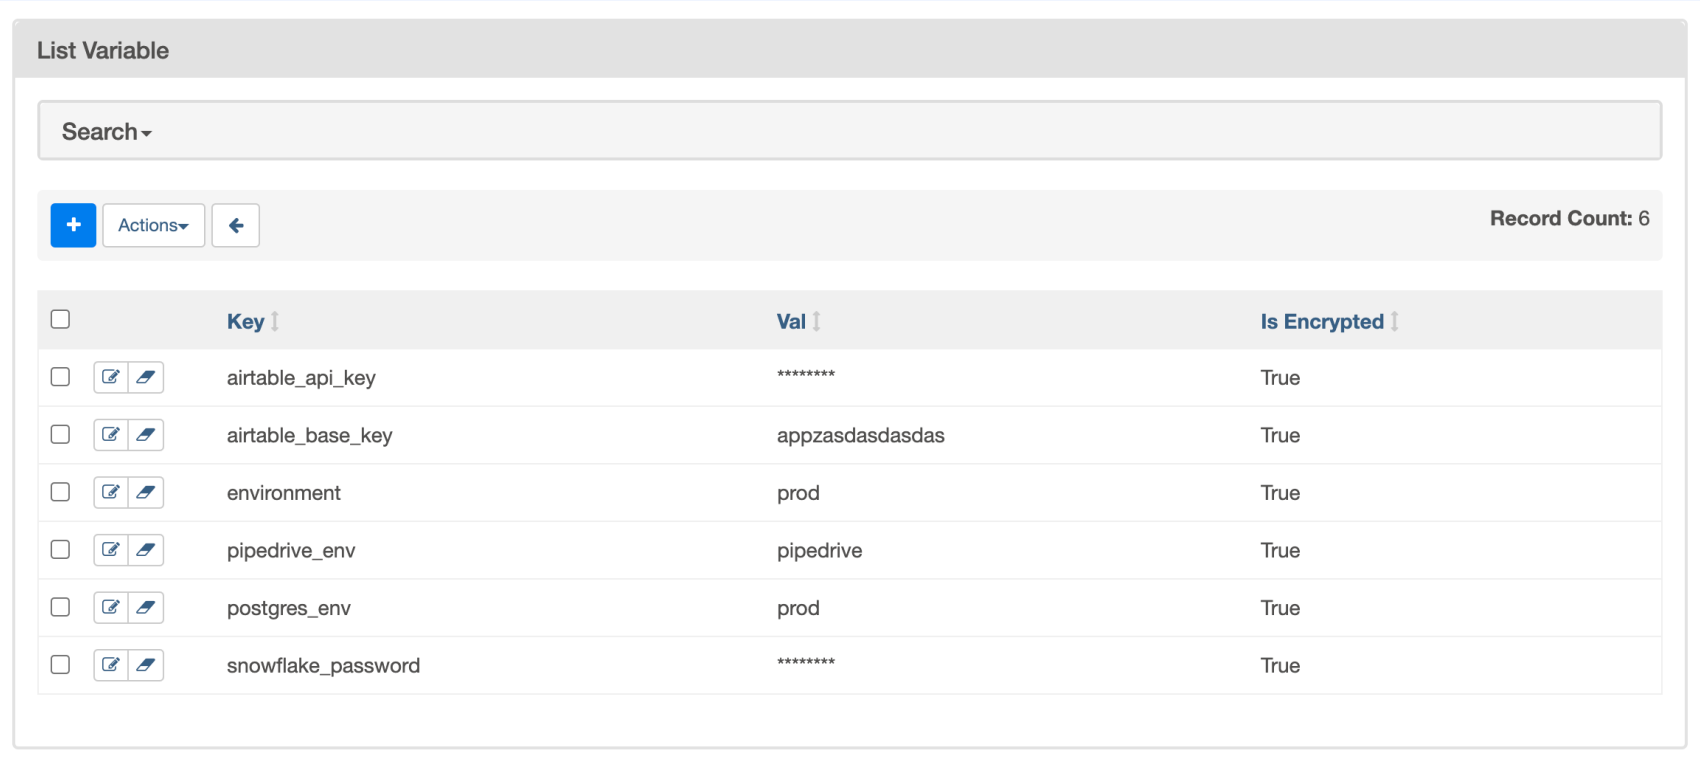

11. Storing Variables in Environment Variables

```
export AIRFLOW_VAR_FOO=BAR

# To use JSON, store them as JSON strings
export AIRFLOW_VAR_FOO_BAZ='{"hello":"world"}’

from airflow.models import Variable

foo = Variable.get("foo")
foo_json = Variable.get("foo_baz", deserialize_json=True)
```

11. Securing Variables

- Fernet
- Secrets Backend
  - Amazon
    - SecretsManagerBackend
    - SystemsManagerParameterStoreBackend
  - Google
    - CloudSecretManagerBackend
  - Hashicorp
    - VaultBackend
  - Microsoft Azure
    - AzureKeyVaultBackend

Chapter 2.

12. Params

12. Params

- Provide runtime configuration to tasks

12. Params vs Variables

- Params:
  - Specific DAG or Task's runtime configuration
  - Charateristic:
    - Provide different values each time you run a DAG or Task
    - Validation is possible using JSON Schema.
  - Store: In DAG or Task definition
- Variables:
  - Global key/value
  - Characteristic: Create, modify and delete via web UI, code, CLI
  - Store: Metadata db (can be secured)

12. DAG-level Params

```
from airflow import DAG
from airflow.models.param import Param

with DAG(
    "the_dag",
    params={
        "x": Param(5, type="integer", minimum=3),
        "my_int_param": 6
    },
):
```

12. Task-level Params

```
def print_my_int_param(params):
  print(params.my_int_param)

PythonOperator(
    task_id="print_my_int_param",
    params={"my_int_param": 10},
    python_callable=print_my_int_param,
)
```

12. Referencing Params in a Task

```
PythonOperator(
    task_id="from_template",
    op_args=[
        "{{ params.my_int_param + 10 }}",
    ],
    python_callable=(
        lambda my_int_param: print(my_int_param)
    ),
)
```

12. Referencing Params in a Task

```
with DAG(
    "the_dag",
    params={"my_int_param": Param(5, type="integer", minimum=3)},
    render_template_as_native_obj=True
):
```

12. Referencing Params in a Task

prints <class 'str'> by default

prints <class 'int'> if render_template_as_native_obj=True

```
PythonOperator(
    task_id="template_type",
    op_args=[
        "{{ params.my_int_param }}",
    ],
    python_callable=(
        lambda my_int_param: print(type(my_int_param))
    ),
)
```

12. Referencing Params in a Task

```
def print_x(**context):
    print(context["params"]["my_int_param"])


PythonOperator(
    task_id="print_my_int_param",
    python_callable=print_my_int_param,
)
```

12. JSON Schema Validation

```
with DAG(
    "my_dag",
    params={
        # a required param which can be of multiple types
        # a param must have a default value
        "multi_type_param": Param(5, type=["null", "number", "string"]),

        # an enum param, must be one of three values
        "enum_param": Param("foo", enum=["foo", "bar", 42]),
    }
):
```

12. Trigger UI Form

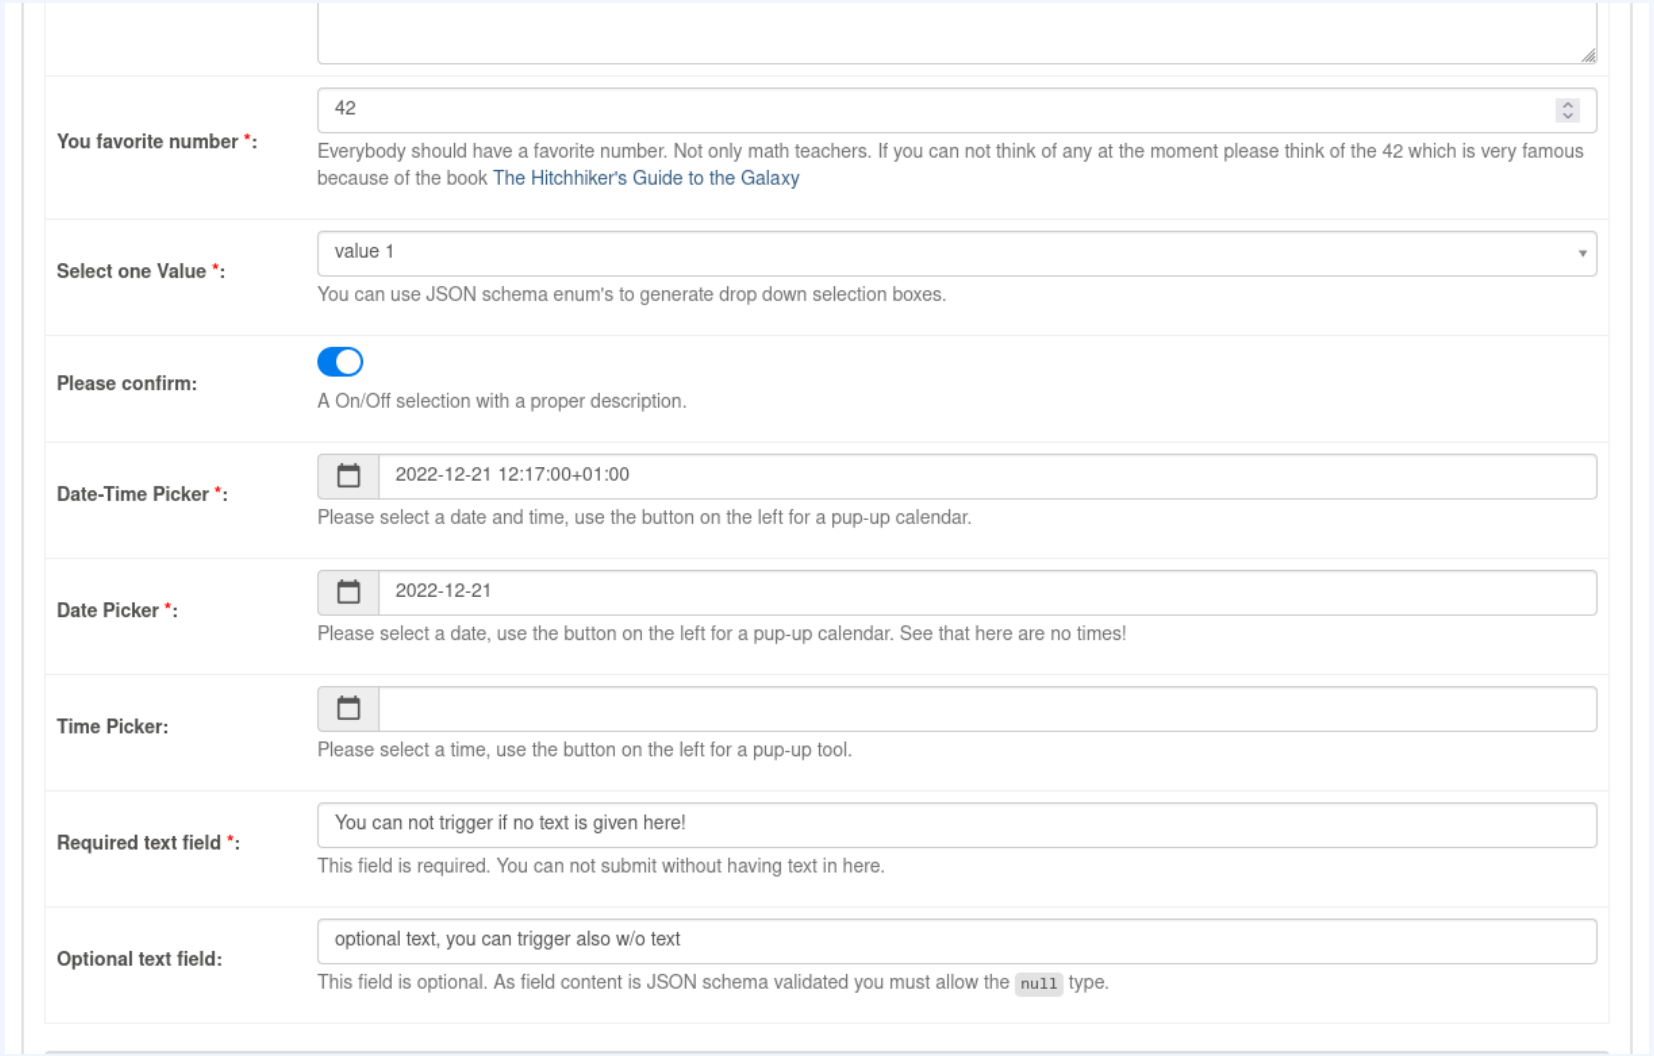

Chapter 2.

13. Configuration

13. Celery

- pool
  - Celery Pool 구현

    가능한 선택사항은 prefork (기본값), eventlet, gevent 또는 solo
  - 타입: 문자열
  - 기본값: prefork
  - 환경 변수: AIRFLOW__CELERY__POOL

13. Celery

- worker_concurrency
  - Worker가 가져올 task instance의 수를 정의
  - 타입: 문자열
  - 기본값: 16
  - 환경 변수: AIRFLOW__CELERY__WORKER_CONCURRENCY

13. Celery

- worker_autoscale
  - Worker를 시작할 때 사용될 최대 및 최소 동시성
  - 타입: 문자열
  - 기본값: None
  - 환경 변수: AIRFLOW__CELERY__WORKER_AUTOSCALE
  - 예제: 16,12

13. Celery

- worker_prefetch_multiplier
  - 성능을 향상시키기 위해 Worker가 미리 가져오는 task의 수를 늘리는 데 사용
  - 타입: 정수
  - 기본값: 1
  - 환경 변수: AIRFLOW__CELERY__WORKER_PREFETCH_MULTIPLIER

13. Core

- default_pool_task_slot_count
  - Airflow 2.2.0 버전에서 새롭게 추가
  - 기본 pool에 대한 task slot 수
  - default_pool이 이미 생성된 기존 배포에서는 영향을 주지 않음
    - 기존 배포에 대해서는 웹서버, API 또는 CLI를 사용하여 변경
  - 타입: 문자열
  - 기본값: 128
  - 환경 변수: AIRFLOW__CORE__DEFAULT_POOL_TASK_SLOT_COUNT

13. Core

- max_active_runs_per_dag
  - DAG당 active DAG Run의 최대 수
  - 타입: 문자열
  - 기본값: 16
  - 환경 변수: AIRFLOW__CORE__MAX_ACTIVE_RUNS_PER_DAG

13. Core

- max_active_tasks_per_dag
  - Airflow 2.2.0 버전에서 새롭게 추가
  - 각 DAG에서 동시에 실행할 수 있는 task instance의 최대 수
    - 새로운 DAG가 클러스터의 모든 Executor slot을 차지하는 걸 방지
  - 타입: 문자열
  - 기본값: 16
  - 환경 변수: AIRFLOW__CORE__MAX_ACTIVE_TASKS_PER_DAG

13. Core

- parallelism
  - Scheduler당 동시에 실행할 수 있는 task instance의 최대 수
    - 해당 값과 클러스터의 Scheduler 수를 곱한 값이 ”running” 상태를 가진 task instance의 최대 수
  - 타입: 문자열
  - 기본값: 32
  - 환경 변수: AIRFLOW__CORE__PARALLELISM

13. Core

- task_runner
  - 하위 프로세스에서 task instance를 실행하는 데 사용할 클래스
    - StandardTaskRunner
    - CgroupTaskRunner
    - Custom TaskRunner
  - 타입: 문자열
  - 기본값: StandardTaskRunner
  - 환경 변수: AIRFLOW__CORE__TASK_RUNNER

13. Scheduler

- parsing_processes
  - Scheduler는 DAG를 파싱하기 위해 여러 프로세스를 동시에 실행
    - 실행될 프로세스의 수를 정의
  - 타입: 문자열
  - 기본값: 2
  - 환경 변수: AIRFLOW__SCHEDULER__PARSING_PROCESSES

13. Summary

- Airflow 구성:
  - Executor 선택
    - executor 설정을 통해 선택
  - 병렬성 및 동시성
    - parallelism: 전체 시스템에서 동시에 실행할 수 있는 task의 최대 수
    - max_active_tasks_per_dag: 각 DAG에서 동시에 실행할 수 있는 task의 최대 수
    - worker_concurrency: 각 Worker에서 동시에 실행할 수 있는 task의 최대 수
- 데이터베이스
  - 풀링
    - sql_alchemy_pool_size 및 sql_alchemy_max_overflow 설정을 통해 조절